

Demo Runbook & Q&A Prep

Pre-session checklist · demo choreography · faculty Q&A answers

Before the session (checklist)

- ☐ Submit the deck + **MVP URL (mandatory): <https://repmint.vercel.app>** + GitHub URL: <https://github.com/simplyYK/repmint> — *before the slot, hard deadline*.
- ☐ Demo machine: signed in at repmint.vercel.app as the demo account, Coach + Workouts + Community tabs pre-opened. Volume UP (voice cue must be heard).
- ☐ Camera permission already granted to the browser on the demo machine.
- ☐ **Record the backup video the evening before** (screen + phone recording of the exact 4 demo beats, ~3 min). Open it muted in a second tab.
- ☐ Phone hotspot ready in case venue wifi blocks WebRTC (realtime voice).
- ☐ One dry run in the actual room if possible — check projector contrast (deck is dark-theme; bump screen brightness).
- ☐ Faculty may interact with the MVP: have the QR slide (7 or 12) up during Q&A; the app is public and sign-up works.

Demo choreography

#	Beat	Where	Time	Wow moment
1	Camera coach	/train	90s	Half-rep rejected + voice cue out loud
2	AI day workout	/workouts → Create	45s	Full session saved in seconds, plan untouched
3	Grounded chat	/coach	45s	Reply cites real session data (July 7, ROM 91–93%)
4	Competitions	/social	30s	Live leaderboard between friends

Failure protocol: anything hangs >10s → switch to the backup video tab and keep narrating the same beats. One light line ("the demo gods want the video version") — never apologize twice, never debug on stage.

Likely faculty questions — strong answers

"Is the AI actually necessary here?"

The core feature is computer vision — without it nothing counts your reps, so a form or dashboard cannot replace it. The LLM layer is also load-bearing: plans

must adapt to goal, equipment, time and injuries; that's generation, not lookup. Deterministic logic (rep thresholds, scoring, clamps) stays deterministic.

"How do you prevent hallucinations?"

Structured outputs validated against a 115-exercise allowlist; invalid JSON is re-asked once, then rejected; every numeric field is clamped server-side. The chat coach is grounded in the user's actual session rows and declines off-topic/medical questions. Where evidence is thin, it's instructed to say so.

"What about privacy — you're filming people?"

We never receive the video. Pose estimation runs in the browser on-device; only landmark-derived numbers (reps, scores) sync. That's architecture, not a policy promise. The database enforces row-level security on every table, and the friends leaderboard is a guarded aggregate function — raw sessions stay private even from friends.

"How accurate is the rep counting?"

Tiered honestly: 58 exercises get full form coaching, 18 get rep counting, 39 are timer/manual — we never pretend the camera can track what it can't. Range-of-motion gates reject partial reps; a One-Euro filter smooths landmark noise.

"Why would anyone pay?"

The free tier proves the magic in 30 seconds at zero marginal cost to us (vision runs on the user's device). Premium unlocks the brain and the voice — about €0.05 of AI cost per active training day against €9.99/month.

"What's defensible here?"

Execution speed and the data loop. Every session generates labeled form data tied to outcomes; the coach's long-term memory makes switching costs real; and competitions make the product social — you can't take your leaderboard with you.

"Realtime voice — doesn't it talk over you or lag?"

One WebRTC session per workout (~300ms latency). It physically cannot self-trigger: no microphone attached, server-side turn detection disabled, it speaks only cues we queue, and it closes itself after the final line. Fallback chain: realtime → HTTP TTS → on-device speech — the workout never goes silent.

"What did you validate with users?"

Three personas from interviews drove scope (duration estimates, short

versions, privacy-by-design, streak saver). The canvas principle guided us: smallest working version that proves value — deployed and testable by you now.

"What breaks at 10,000 users?"

Nothing architectural: static frontend on a CDN, vision cost scales with users' own devices, Postgres + edge functions scale managed. The bottleneck would be LLM spend — bounded per action and cached exercise bank keeps prompts lean.

"Why isn't feature X in the MVP?"

Deliberate scope: we shipped the loop that proves value end-to-end (see → coach → plan → remember → compete). The canvas's own rule: a smaller working MVP with evidence beats an ambitious fake demo.