

# Demo Runbook & Q&A Prep

Pre-session checklist · demo choreography · faculty Q&A answers

## Before the session (checklist)

- ☐ Submit the deck + **MVP URL (mandatory): <https://repmint.vercel.app>** + GitHub URL: <https://github.com/simplyYK/repmint> — *before the slot, hard deadline*.
- ☐ Demo machine: signed in at [repmint.vercel.app](https://repmint.vercel.app) as the demo account, Coach + Workouts + Community tabs pre-opened. Volume UP (voice cue must be heard).
- ☐ Camera permission already granted to the browser on the demo machine.
- ☐ **Record the backup video the evening before** (screen + phone recording of the exact 4 demo beats, ~3 min). Open it muted in a second tab.
- ☐ Phone hotspot ready in case venue wifi blocks WebRTC (realtime voice).
- ☐ One dry run in the actual room if possible — check projector contrast (deck is dark-theme; bump screen brightness).
- ☐ Faculty may interact with the MVP: have a QR slide (9 or 14) up during Q&A; the app is public and sign-up works.

## Demo choreography

| # | Beat           | Where              | Time | Wow moment                                         |
|---|----------------|--------------------|------|----------------------------------------------------|
| 1 | Camera coach   | /train             | 90s  | Half-rep rejected + voice cue out loud             |
| 2 | AI day workout | /workouts → Create | 45s  | Full session saved in seconds, plan untouched      |
| 3 | Grounded chat  | /coach             | 45s  | Reply cites real session data (July 7, ROM 91–93%) |
| 4 | Competitions   | /social            | 30s  | Live leaderboard between friends                   |

**Failure protocol:** anything hangs >10s → switch to the backup video tab and keep narrating the same beats. One light line ("the demo gods want the video version") — never apologize twice, never debug on stage.

## Likely faculty questions — strong answers

**"Is the AI actually necessary here?"**

The core feature is computer vision — without it nothing counts your reps, so a form or dashboard cannot replace it. The LLM layer is also load-bearing: plans

must adapt to goal, equipment, time and injuries; that's generation, not lookup. Deterministic logic (rep thresholds, scoring, clamps) stays deterministic.

### **"How do you prevent hallucinations?"**

Structured outputs validated against a 115-exercise allowlist; invalid JSON is re-asked once, then rejected; every numeric field is clamped server-side. The chat coach is grounded in the user's actual session rows and declines off-topic/medical questions. Where evidence is thin, it's instructed to say so.

### **"What about privacy — you're filming people?"**

We never receive the video. Pose estimation runs in the browser on-device; only landmark-derived numbers (reps, scores) sync. That's architecture, not a policy promise. The database enforces row-level security on every table, and the friends leaderboard is a guarded aggregate function — raw sessions stay private even from friends.

### **"How accurate is the rep counting?"**

Tiered honestly: 58 exercises get full form coaching, 18 get rep counting, 39 are timer/manual — we never pretend the camera can track what it can't. Range-of-motion gates reject partial reps; a One-Euro filter smooths landmark noise.

### **"Why would anyone pay?"**

The free tier proves the magic in 30 seconds at zero marginal cost to us (vision runs on the user's device). Premium unlocks the brain and the voice — about €0.05 of AI cost per active training day against €9.99/month.

### **"What's defensible here?"**

Execution speed and the data loop. Every session generates labeled form data tied to outcomes; the coach's long-term memory makes switching costs real; and competitions make the product social — you can't take your leaderboard with you.

### **"Realtime voice — doesn't it talk over you or lag?"**

One WebRTC session per workout (~300ms latency). It physically cannot self-trigger: no microphone attached, server-side turn detection disabled, it speaks only cues we queue, and it closes itself after the final line. Fallback chain: realtime → HTTP TTS → on-device speech — the workout never goes silent.

### **"What did you validate with users?"**

Three personas from interviews drove scope (duration estimates, short

versions, privacy-by-design, streak saver). The canvas principle guided us: smallest working version that proves value — deployed and testable by you now.

**"What breaks at 10,000 users?"**

Nothing architectural: static frontend on a CDN, vision cost scales with users' own devices, Postgres + edge functions scale managed. The bottleneck would be LLM spend — bounded per action and cached exercise bank keeps prompts lean.

**"Why isn't feature X in the MVP?"**

Deliberate scope: we shipped the loop that proves value end-to-end (see → coach → plan → remember → compete). The canvas's own rule: a smaller working MVP with evidence beats an ambitious fake demo.