

Presentation Script — 4 Speakers

AI MVP Challenge · 10 minutes incl. live demo · repmint.vercel.app

Total: 9:45 + buffer. Deck: `RepMint-MVP-Deck.pptx` (14 slides). Demo: repmint.vercel.app

Faculty rule: reading from paper or screen is penalized. These are *beats to internalize*, not lines to read. Each speaker owns their beats; rehearse the handoff sentences out loud — they're the glue.

SPEAKER 1 — The Problem (0:00 – 2:00) · Slides 1–4

[Slide 1 — Title] (35s)

Quick show of hands — who has ever propped up their phone at the gym to film a set, just to check their own form afterwards?

That awkward video exists because when you train alone, nobody is watching. We built the coach that watches. We're Team RepMint — [first names, one breath] — and RepMint is a personal trainer that lives in your camera.

[Slide 2 — Problem] (45s)

Here's the gap. A personal trainer costs fifty to eighty euros a session — that's not a habit, that's a luxury. So four out of five gym-goers train with zero feedback: no one counts an honest rep, no one catches a knee caving in, no one tells you when to stop.

And every fitness app on your phone right now has the same blind spot — it logs whatever you *type*. It has no idea what you actually *did*. The moment of pain is mid-set, alone: "was that rep even right?"

[Slide 4 — NABC] (40s)

The need: real-time form guidance for people who will never pay for a trainer. Our approach: a 33-point pose model that runs entirely on your phone, in the browser — no app store, no hardware. The benefit: a coach that sees, speaks, plans and remembers, for the price of a coffee — and your video never leaves the device.

Competition? Freeletics and Fitbod plan workouts but they can't *see* you. Smart mirrors see you — for two thousand euros of hardware. Trainers see you but don't scale. We are the only camera-native, software-only coach.

[Slide 3 — "Nobody is watching."] (10s beat)

(Say the line. Let it sit for two full seconds — the empty gym does the work. Then click.) Nobody is watching. And form is the first thing that breaks.

HANDOFF → "So who is this for? [Speaker 2] met them."

SPEAKER 2 — The People & The Product (2:00 – 3:45) · Slides 5–7

[Slide 5 — Personas] (35s)

Three people from our user research. Mateo, 26, trains between classes — he abandons any app that makes him take more than two decisions. Priya, 22, dorm room, no equipment — she'll never build a workout from a blank screen, and she cares deeply that camera video stays private. Lukas, 29, ex-athlete — what keeps him going is the streak and his friends seeing it.

Every feature you're about to see exists because one of these three needed it.

[Slide 6 — Jump Start] (30s)

We didn't spend three weeks on slides. Week one: the camera engine — pose tracking, honest rep counting, a 115-exercise library. Week two: the brain — AI plans, a coach with long-term memory, full training history. Week three: the feel — a realtime voice, a community, and hardening. Forty-five-plus production deployments later, it's live at repmint-dot-vercel-dot-app.

[Slide 7 — Solution] (40s)

One coach, five superpowers. It **sees**: only full range-of-motion reps count — half reps don't fool it. It **speaks**: a natural voice mid-set, one cue at a time. It **thinks**: weekly plans, day workouts on demand, answers grounded in your real data. It **connects**: friends and week-long competitions. And it **protects**: pose runs on-device; video is never uploaded, ever.

And to answer the canvas question directly — is the AI necessary or decorative? Without the vision model there is no product. A form can't count your reps.

HANDOFF → Click to slide 8 ("Enough slides.") AS you say: "Enough slides. [Speaker 3] — show them." Speaker 3 walks to the demo machine over it.

SPEAKER 3 — Live Demo (3:45 – 7:15) · Slide 9 stays up, then the app

*(Slide 8 is your entrance music; click to slide 9 — run-of-show + QR while you switch to the browser.

Phone/laptop already signed in. Backup video open in a second tab, muted.)*

Beat 1 — Train (90s, the wow):

- Open Train → pick the day's workout → camera gate → prop the phone.
- Step back, raise a hand → "3·2·1" → do 4–5 squats live on stage.
- Narrate only what the room can see: *"It's counting — but watch, I'll cut this one short... no count. Only honest reps."*
- Let ONE voice cue play out loud. *"That's the realtime coach — about 300 milliseconds from movement to voice."*
- Finish the set → point at per-rep quality + fatigue read.

Beat 2 — AI builds today's workout (45s):

- Workouts → Create workout → "Ask your AI coach."
- Type: *"chest and triceps, 30 minutes, feeling strong"* → Build.
- While it generates: *"This builds ONE session for today without touching my weekly plan — different job than the planner."*
- Show the saved workout with exercises.

Beat 3 — A coach that remembers (45s):

- Coach → open "Fixing my squat depth" → scroll the history.
- Ask live: "*did the pause trick work?*" → the reply cites the actual July 7 session, weights and ROM scores. *"It's not chatting — it's reading my training data."*
- Point at markdown rendering, regenerate button, stop button — *"full chat controls."*

Beat 4 — Accountability (30s):

- Community → live competition leaderboard → *"week-long metric races between friends. This is our retention loop AND our growth loop."*

FALLBACK: any hiccup > 10 seconds → switch tabs to the backup video and keep narrating over it with the same beats. Do not apologize twice.

HANDOFF → "It works. [Speaker 4] — why it's a business."

SPEAKER 4 — Business, Architecture, Close (7:15 – 9:45) · Slides 10–14

[Slide 10 — Business value] (50s)

The business writes itself from the architecture. Free tier: the camera coach — because it shows the magic in thirty seconds and costs us nothing; pose runs on the *user's* device. Premium, €9.99 a month: the AI brain and the realtime voice — our marginal AI cost is about five cents per active training day. That's software margins on a personal-training problem. And the community layer means the product invites the next user for us — competitions literally require friends.

[Slide 11 — Architecture] (40s)

One idea to remember: heavy where it matters, light everywhere else. The heaviest AI — vision — runs in the browser. That single decision makes RepMint private, free to scale, and instant. The cloud does only what must be shared: Postgres with row-level security on every table, six edge functions, and the language models — OpenRouter with a Gemini fallback, and OpenAI's Realtime API for the voice over WebRTC.

[Slide 12 — Reliability] (30s)

We built for trust, not just for demo day. Two examples: generated plans are validated against a 115-exercise allowlist — the model *cannot* invent an exercise. And the voice agent physically cannot ramble: no microphone, turn detection off, it speaks only the cues we send, and the session closes the moment the workout ends.

[Slide 13 — Future & conclusions] (30s)

Next: wearables, form-video library, a workout marketplace, B2B pilots with gyms and university sport programs. But the conclusion is on the right: the full loop — see, coach, plan, remember, compete — works today, deployed, for anyone in this room. Your own canvas says it best: a smaller working MVP with evidence beats an ambitious fake demo.

[Slide 14 — Close] *(10s)*

We're Team RepMint. Stop guessing, start counting — scan it and train. Questions welcome.

(All four stand for Q&A.)