

Presentation Script — 5 Speakers

AI MVP Challenge · 10 minutes incl. live demo · repmint.vercel.app

Total: 9:45 + buffer. Deck: `RepMint-MVP-Deck.pptx` (12 slides). Demo: repmint.vercel.app

Same beats as the 4-speaker version, with the demo split across two presenters so everyone owns real content (faculty may direct questions at non-presenters — in this version there are none).

SPEAKER 1 — Hook, Problem, NABC (0:00 – 1:50) · Slides 1–3

Same as 4-speaker Speaker 1, trimmed ~10s per slide:

Hook: who has filmed their own set to check form? That video exists because nobody is watching. We're Team RepMint — [names] — and RepMint is a personal trainer that lives in your camera.

Problem: trainers cost €50–80 a session; four in five gym-goers train with zero feedback; and every app is blind — it logs what you type, not what you do.

NABC: Need — real-time form guidance without a trainer's price tag. Approach — 33-point pose AI fully on-device in the browser. Benefits — a coach that sees, speaks, plans, remembers; video never leaves the phone. Competition — apps that can't see you, mirrors that cost €2,000, trainers that don't scale. We're the only camera-native software coach.

HANDOFF → "Who needs it? [Speaker 2] met them."

SPEAKER 2 — Personas, Jump Start, Solution (1:50 – 3:20) · Slides 4–6

Same beats as 4-speaker Speaker 2:

Mateo optimizes minutes between classes. Priya trains in a dorm and worries about camera privacy. Lukas lives for the streak and his friends' respect.

Three weeks: week one the camera engine, week two the AI brain, week three voice, community and polish. 45+ deployments, live right now — every exercise page now carries our own AI-illustrated muscle maps.

Five superpowers: it sees (honest reps only), speaks (realtime voice), thinks (plans + grounded chat), connects (competitions), protects (video never uploaded). Is the AI decorative? Without the camera model there is no product.

HANDOFF → "[Speaker 3] — show them the coach."

SPEAKER 3 — Demo Part 1: The Camera Coach (3:20 – 5:20) · Slide 7 → app

Beat 1 — Train flow (100s):

- Train → today's workout → camera gate → prop the phone → step back.
- Raise a hand → countdown → 4–5 live squats.
- *"Watch the count — and watch this half rep... rejected. Only honest reps."*
- Let one voice cue play: *"realtime voice, ~300ms from movement to speech."*
- End of set: per-rep quality, time under tension, fatigue read.
- End of workout beat: *"and when the session ends, the coach says its last line and goes quiet — the voice session closes itself."*

Beat 2 — AI day workout (20s):

- Create workout → "Ask your AI coach" → *"chest and triceps, 30 minutes"* → it saves a full session without touching the weekly plan.

HANDOFF → "That's the body. [Speaker 4] — the brain and the friends."

SPEAKER 4 — Demo Part 2: Memory & Community (5:20 – 7:00) · app

Beat 3 — Grounded coach chat (50s):

- Coach → "Fixing my squat depth" conversation → ask live: *"did the pause trick work?"*
- Reply cites the real July 7 session — weights, ROM scores. *"Not a chatbot — it reads my training data. And it remembers my tweaky shoulder across every conversation."*
- Flash the chat controls: markdown, stop, regenerate.

Beat 4 — Competitions (40s):

- Community → leaderboard of this week's competition → join/create flow.
- *"Week-long races on reps, sets, sessions or minutes — between friends only. Accountability is the feature that keeps people training; it's also the feature that invites the next user."*

FALLBACK (both demo speakers): hiccup > 10s → backup video tab, same narration, no second apology.

HANDOFF → "It works. [Speaker 5] — why it wins."

SPEAKER 5 — Business, Architecture, Reliability, Close (7:00 – 9:45) · Slides 8–12

Same as 4-speaker Speaker 4, at full length:

Business: free camera coach (zero marginal cost — vision runs on the user's device), €9.99 premium for the AI brain and voice (~€0.05 per active day → software margins), community as the built-in growth loop.

Architecture: heavy where it matters, light everywhere else — vision in the browser; the cloud only for what must be shared: Postgres + RLS on every table, six edge functions, OpenRouter→Gemini fallback, OpenAI Realtime voice over WebRTC.

Reliability: plans validated against a 115-exercise allowlist — the model cannot invent movements; the voice agent cannot ramble — no mic, no turn detection, cue-driven only, session closes when the workout ends.

Future: wearables, form videos, marketplace, B2B gym pilots.

Close: the full loop works today, deployed, for anyone in this room. A smaller working MVP with evidence beats an ambitious fake demo. We're Team RepMint — scan it and train.

(All five stand for Q&A.)